Serverless Computing / Microservices / Cloud Native

An Introduction to OpenWhisk, IBM's Open Source Serverless Computing Platform

6 Jun 2017 1:15am, by Jonathan Mitt

Serverless computing, sometimes referred to as serverless computing, emerge with Lambda in 2014, no one expected it to become an alternative computing model to virtual machines and containers. But since its launch, Lambda has evolved to become a key component of AWS compute services portfolio. Deeper integration with mainstream AWS services, support for popular languages, streamlined deployment model all make Lambda a viable microservices platform for AWS customers.

Though AWS Lambda caught the attention of the industry and developers, there is a general concern about lock-in. An application developed for Lambda cannot run outside the context of AWS. For customers considering an alternative that can run within the datacenter and cloud, OpenWhisk from IBM can be a potential option.

Announced in February 2016, OpenWhisk is an event-driven computing framework from IBM. It is integrated with Bluemix, IBM's Cloud Foundry-based PaaS and is also available as an open source software that can be deployed behind the firewall. Adobe and IBM jointly submitted OpenWhisk to Apache, which has now become an incubation project. It is one of the first open source serverless platforms in the market. In December 2016, IBM announced the general availability of Bluemix OpenWhisk, the hosted version in the public cloud.

OpenWhisk Concepts and Terminology
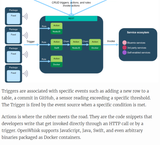
Like other serverless frameworks, OpenWhisk invokes code snippets when a specific event occurs. There are three key concepts in OpenWhisk: Triggers,